

Causal Inference with `group_by` and `summarise`


Lucy D'Agostino McGowan

Wake Forest University

2022-07-23 (updated: 2022-10-16)


Observational Studies

Goal: To answer a research question



Observational Studies

Goal: To answer a research question




Observational Studies

Randomized Controlled Trial



Observational Studies

Randomized Controlled Trial




Observational Studies








Confounding



Confounding



One binary confounder

Simulation

```
n <- 1000
sim <- tibble(
  confounder = rbinom(n, 1, 0.5),
  p_exposure = case_when(
    confounder == 1 ~ 0.75,
    confounder == 0 ~ 0.25
  ),
  exposure = rbinom(n, 1, p_exposure),
  outcome = confounder + rnorm(n)
)
```

```
## # A tibble: 1,000 × 3
##   confounder exposure outcome
##   <int>      <int>    <dbl>
## 1         0         0     1.13
## 2         0         0     1.11
## 3         1         1     0.129
## 4         1         0     1.21
## 5         0         0     0.0694
## 6         1         1    -0.663
## 7         1         1     1.81
## 8         1         1    -0.912
## 9         1         0    -0.247
## 10        0         0     0.998
## # ... with 990 more rows
```

Simulation

```
n <- 1000
sim <- tibble(
  confounder = rbinom(n, 1, 0.5),
  p_exposure = case_when(
    confounder == 1 ~ 0.75,
    confounder == 0 ~ 0.25
  ),
  exposure = rbinom(n, 1, p_exposure),
  outcome = confounder + rnorm(n)
)
```

```
## # A tibble: 1,000 × 3
##   confounder exposure outcome
##   <int>      <int>    <dbl>
## 1         0         0     1.13
## 2         0         0     1.11
## 3         1         1     0.129
## 4         1         0     1.21
## 5         0         0     0.0694
## 6         1         1    -0.663
## 7         1         1     1.81
## 8         1         1    -0.912
## 9         1         0    -0.247
## 10        0         0     0.998
## # ... with 990 more rows
```

Simulation

```
n <- 1000
sim <- tibble(
  confounder = rbinom(n, 1, 0.5),
  p_exposure = case_when(
    confounder == 1 ~ 0.75,
    confounder == 0 ~ 0.25
  ),
  exposure = rbinom(n, 1, p_exposure),
  outcome = confounder + rnorm(n)
)
```

```
## # A tibble: 1,000 × 3
##   confounder exposure outcome
##   <int>      <int>    <dbl>
## 1         0         0    1.13
## 2         0         0    1.11
## 3         1         1    0.129
## 4         1         0    1.21
## 5         0         0    0.0694
## 6         1         1   -0.663
## 7         1         1    1.81
## 8         1         1   -0.912
## 9         1         0   -0.247
## 10        0         0    0.998
## # ... with 990 more rows
```

Simulation

```
n <- 1000
sim <- tibble(
  confounder = rbinom(n, 1, 0.5),
  p_exposure = case_when(
    confounder == 1 ~ 0.75,
    confounder == 0 ~ 0.25
  ),
  exposure = rbinom(n, 1, p_exposure),
  outcome = confounder + rnorm(n)
)
```

```
## # A tibble: 1,000 × 3
##   confounder exposure outcome
##   <int>      <int>    <dbl>
## 1         0         0     1.13
## 2         0         0     1.11
## 3         1         1     0.129
## 4         1         0     1.21
## 5         0         0     0.0694
## 6         1         1    -0.663
## 7         1         1     1.81
## 8         1         1    -0.912
## 9         1         0    -0.247
## 10        0         0     0.998
## # ... with 990 more rows
```

Simulation

```
n <- 1000
sim <- tibble(
  confounder = rbinom(n, 1, 0.5),
  p_exposure = case_when(
    confounder == 1 ~ 0.75,
    confounder == 0 ~ 0.25
  ),
  exposure = rbinom(n, 1, p_exposure),
  outcome = confounder + rnorm(n)
)
```

```
## # A tibble: 1,000 × 3
##   confounder exposure outcome
##   <int>      <int>    <dbl>
## 1         0         0     1.13
## 2         0         0     1.11
## 3         1         1     0.129
## 4         1         0     1.21
## 5         0         0     0.0694
## 6         1         1    -0.663
## 7         1         1     1.81
## 8         1         1    -0.912
## 9         1         0    -0.247
## 10        0         0     0.998
## # ... with 990 more rows
```


Simulation

```
lm(outcome ~ exposure, data = sim)
```

```
##
```

```
## Call:
```

```
## lm(formula = outcome ~ exposure, data = sim)
```

```
##
```

```
## Coefficients:
```

```
## (Intercept)      exposure
```

```
##      0.2688      0.4070
```

Simulation

```
sim %>%  
  group_by(exposure) %>%  
  summarise(avg_y = mean(outcome))
```

```
## # A tibble: 2 × 2  
##   exposure avg_y  
##   <int> <dbl>  
## 1      0 0.269  
## 2      1 0.676
```

Simulation

```
sim %>%  
  group_by(exposure) %>%  
  summarise(avg_y = mean(outcome))
```

```
## # A tibble: 2 × 2  
##   exposure avg_y  
##   <int> <dbl>  
## 1      0 0.269  
## 2      1 0.676
```

Simulation

```
sim %>%  
  group_by(exposure) %>%  
  summarise(avg_y = mean(outcome))
```

```
## # A tibble: 2 × 2  
##   exposure avg_y  
##   <int> <dbl>  
## 1      0 0.269  
## 2      1 0.676
```

Simulation

```
sim %>%  
  group_by(exposure) %>%  
  summarise(avg_y = mean(outcome)) %>%  
  pivot_wider(names_from = exposure,  
              values_from = avg_y,  
              names_prefix = "x_") %>%  
  summarise(estimate = x_1 - x_0)
```

```
## # A tibble: 1 × 1  
##   estimate  
##   <dbl>  
## 1      0.407
```

Simulation

```
sim %>%  
  group_by(confounder, exposure) %>%  
  summarise(avg_y = mean(outcome))
```

```
## # A tibble: 4 × 3  
## # Groups:   confounder [2]  
##   confounder exposure    avg_y  
##      <int>      <int>    <dbl>  
## 1         0         0 -0.00907  
## 2         0         1 -0.0166  
## 3         1         0  1.09  
## 4         1         1  0.936
```

Simulation

```
sim %>%  
  group_by(confounder, exposure) %>%  
  summarise(avg_y = mean(outcome)) %>%  
  pivot_wider(names_from = exposure,  
              values_from = avg_y,  
              names_prefix = "x_") %>%  
  summarise(estimate = x_1 - x_0)
```

```
## # A tibble: 2 × 2  
##   confounder estimate  
##   <int>     <dbl>  
## 1         0 -0.00752  
## 2         1 -0.151
```



Two binary confounders

Simulation

```
n <- 1000
sim2 <- tibble(
  confounder_1 = rbinom(n, 1, 0.5),
  confounder_2 = rbinom(n, 1, 0.5),

  p_exposure = case_when(
    confounder_1 == 1 & confounder_2 == 1 ~ 0.5,
    confounder_1 == 0 & confounder_2 == 1 ~ 0.5,
    confounder_1 == 1 & confounder_2 == 0 ~ 0.5,
    confounder_1 == 0 & confounder_2 == 0 ~ 0.5
  ),
  exposure = rbinom(n, 1, p_exposure),
  outcome = confounder_1 + confounder_2 + rnorm(n)
)
```

```
## # A tibble: 1,000 × 4
##   confounder_1 confounder_2 exposure outcome
##   <int>         <int>     <int>    <dbl>
## 1           0           0         0    0.521
## 2           1           0         0    1.38
## 3           0           0         0   -0.624
## 4           0           1         1    0.427
## 5           1           0         1    1.31
## 6           0           0         0   -0.707
## 7           1           1         1    2.52
## 8           1           0         0    1.45
## 9           0           0         0   -0.505
## 10          0           1         1    0.793
## # ... with 990 more rows
```

Simulation

```
lm(outcome ~ exposure, data = sim2)
```

```
##
```

```
## Call:
```

```
## lm(formula = outcome ~ exposure, data = sim2)
```

```
##
```

```
## Coefficients:
```

```
## (Intercept)      exposure
```

```
##      0.6395      0.6951
```

Simulation

```
sim2 %>%  
  group_by(confounder_1, confounder_2, exposure) %>%  
  summarise(avg_y = mean(outcome)) %>%  
  pivot_wider(names_from = exposure,  
              values_from = avg_y,  
              names_prefix = "x_") %>%  
  summarise(estimate = x_1 - x_0)
```

```
## # A tibble: 4 × 3  
## # Groups:   confounder_1 [2]  
##   confounder_1 confounder_2 estimate  
##           <int>         <int>     <dbl>  
## 1             0             0 -0.309  
## 2             0             1  0.0466  
## 3             1             0 -0.0271  
## 4             1             1 -0.00275
```

Simulation

```
n <- 100000
sim2 <- tibble(
  confounder_1 = rbinom(n, 1, 0.5),
  confounder_2 = rbinom(n, 1, 0.5),

  p_exposure = case_when(
    confounder_1 == 1 & confounder_2 == 1 ~ 0.5,
    confounder_1 == 0 & confounder_2 == 1 ~ 0.3,
    confounder_1 == 1 & confounder_2 == 0 ~ 0.4,
    confounder_1 == 0 & confounder_2 == 0 ~ 0.2
  ),
  exposure = rbinom(n, 1, p_exposure),
  outcome = confounder_1 + confounder_2 + rnorm(n)
)
```

```
## # A tibble: 100,000 × 4
##   confounder_1 confounder_2 exposure outcome
##   <int>         <int>     <int>    <dbl>
## 1             1           1         1    2.35
## 2             1           1         0    3.71
## 3             0           0         0    2.08
## 4             0           1         1    0.516
## 5             0           0         0   -0.166
## 6             1           1         1    1.58
## 7             0           0         0    0.472
## 8             1           0         0    3.22
## 9             0           1         1    0.929
## 10            0           1         1    1.41
## # ... with 99,990 more rows
```

Simulation

```
lm(outcome ~ exposure, data = sim2)
```

```
##
```

```
## Call:
```

```
## lm(formula = outcome ~ exposure, data = sim2)
```

```
##
```

```
## Coefficients:
```

```
## (Intercept)      exposure
```

```
##      0.6782      0.6561
```

Simulation

```
sim2 %>%  
  group_by(confounder_1, confounder_2, exposure) %>%  
  summarise(avg_y = mean(outcome)) %>%  
  pivot_wider(names_from = exposure,  
              values_from = avg_y,  
              names_prefix = "x_") %>%  
  summarise(estimate = x_1 - x_0)
```

```
## # A tibble: 4 × 3  
## # Groups:   confounder_1 [2]  
##   confounder_1 confounder_2 estimate  
##         <int>         <int>     <dbl>  
## 1           0           0  0.00565  
## 2           0           1  0.0466  
## 3           1           0  0.0185  
## 4           1           1  0.00388
```

Continuous confounder?

Simulation

```
n <- 10000
sim3 <- tibble(
  confounder = rnorm(n),
  p_exposure = exp(confounder) / (1 + exp(confounder)),
  exposure = rbinom(n, 1, p_exposure),
  outcome = confounder + rnorm(n)
)
```

```
## # A tibble: 10,000 × 3
##   confounder exposure outcome
##   <dbl>      <int>    <dbl>
## 1   -0.167         0   -0.560
## 2    0.252         1    0.628
## 3   -0.321         1  -0.608
## 4    0.621         0    1.58
## 5   -0.619         1    0.358
## 6   -0.897         0   -1.95
## 7   -2.01         0   -2.50
## 8    0.296         0   -1.10
## 9   -0.504         1  -0.316
## 10  -0.536         1    1.12
## # ... with 9,990 more rows
```


Simulation

```
lm(outcome ~ exposure, data = sim3)
```

```
##
```

```
## Call:
```

```
## lm(formula = outcome ~ exposure, data = sim3)
```

```
##
```

```
## Coefficients:
```

```
## (Intercept)      exposure
```

```
##      -0.4036        0.8152
```

Simulation

```
sim3 %>%  
  mutate(confounder_q = ntile(confounder, 5)) %>%  
  group_by(confounder_q, exposure) %>%  
  summarise(avg_y = mean(outcome)) %>%  
  pivot_wider(names_from = exposure,  
              values_from = avg_y,  
              names_prefix = "x_") %>%  
  summarise(estimate = x_1 - x_0)
```

```
## # A tibble: 5 × 2  
##   confounder_q estimate  
##   <int>      <dbl>  
## 1         1    0.104  
## 2         2   -0.0293  
## 3         3    0.0201  
## 4         4    0.0674  
## 5         5    0.201
```

What if we could come up with a
summary score of all
confounders?